

# MP4: Page Manager II

Nikhil Nunna

UIN: 525008698

CSCE611: Operating System

## Assigned Tasks

**Part I:** Completed.

**Part II:** Completed.

**Part III:** Completed code but has bug in `is_legitimate`.

## System Design

The goal of this program was to create a Page Manager that could handle virtual memory. The addition of virtual memory makes it possible for the user to dynamically add and remove memory spaces in the address space. The user does this by calling functions such as `new` and `malloc`, they can even deallocate the memory using `delete`. The virtual memory is larger and dynamically allocated so this means that we have to move from direct mapped memory in the Kernel frame pool to the non-direct mapped memory in the Process frame pool.

In our implementation, we focused on 2 things, re-building the structure of the page table so that it could do recursive page-table look-up so it would be compatible with the process frame pool and modifying the page-fault handler so that it would be able to check if locations were valid in a VM-Pool. Our implementation of the page table has 2 levels meaning we have a page directory which is like a page table of page tables. Then each of these sub-page tables contains pages that translate to frames and they can be dynamically allocated as needed. In order to navigate both of these tables through the recursive page-table look-up I created 2 helper functions as recommended by the MP4 handout. This means that whenever I had to modify a PDE or PTE I would just call the helper function which would return a pointer to that memory location. These could then be dereferenced and used to modify what I needed to in the PTE or PDE. To incorporate the VM pools into the page fault handler all I did was check whether or not the address being modified was in a valid region of memory before actually handling the page fault.

## Code Description

The code compiles and runs normally. All the grader needs to do is run 'make', './copykernel.sh' and then 'bochs -f bochsrc.bxrc'. As an alternative to this, I've created test.sh which can be run using './test.sh'. The test.sh shell script just runs make, copykernel and the bochs command for you. There are 2 changes made to the makefile to use the frame pool object file given to us, one is to exclude it from the make clean, and the other is to comment out the line that makes the object file so the one given to us doesn't get overwritten before linking.

**test.sh: code block** : This shell file compiles and runs the project

```
#!/bin/bash
```

```
make
```

```
./copykernel.sh
```

```
bochs -f bochsrc.bxrc
```

**page\_table.c: constructor** : In this function, there have been some changes since the last MP, the biggest change is we are now using the process pool for our whole page table. The second change was made in order to support the first change, for paging to work in non-direct mapped memory I had to support recursive paging. To do this I simply set the last entry in the page directory equal to the head of the page directory.

```
PageTable::PageTable()
```

```
{  
    // Get page directory frame  
    unsigned long page_directory_address = 4096 * process_mem_pool->get_frames(1);  
    page_directory = (unsigned long *) page_directory_address;  
  
    // Get page table frame  
    unsigned long page_table_address = 4096 * process_mem_pool->get_frames(1);  
    unsigned long *page_table = (unsigned long *) page_table_address;  
  
    // Direct map first 4 MB of memory  
    unsigned long address=0;  
    unsigned int i;  
  
    // Loop thru page table and set all address to present for first 4 mb  
    for(i=0; i<1024; i++)  
    {  
        page_table[i] = address | 3;  
        address = address + 4096;  
    }  
  
    // Add page table to page_directory  
    page_directory[0] = (unsigned long) page_table;  
    page_directory[0] = page_directory[0] | 3;  
  
    // Recursive page table look-up  
    page_directory[1023] = (unsigned long) page_directory;  
    page_directory[1023] = page_directory[1023] | 3;  
  
    // Loop thru page directory and set all other PDEs to empty  
    for(i=1; i<1023; i++)  
    {  
        page_directory[i] = 0UL | 2;  
    }  
}
```

```

    Console::puts("Constructed Page-Table-object\n");
}

```

**page\_table.c: register\_pool** : In this function, I add the VM pool to the page tables list of pools. This is done by checking if there is a VMpool in the list, if there isn't the Head and Tail are set to the new pool. If the list is populated then the new VMpool is appended to the end and made the Tail.

```

void PageTable::register_pool(VMPool * _vm_pool)
{
    // check if there is anything in the list
    // if there isn't add to it
    if (Head == nullptr)
    {
        Head = _vm_pool;
        Tail = _vm_pool;
    }
    // since there is something already in the list append to it
    else
    {
        Tail->next = _vm_pool;
        _vm_pool->prev = Tail;
        Tail = _vm_pool;
    }
    Console::puts("registered VM-pool\n");
}

```

**page\_table.c: free\_page** : In this function, we enable paging by flipping a bit in the cr0 register and setting our object internal data bool "paging-enabled" to true.

```

void PageTable::free_page(unsigned long _page-no) {
    // get the PTE so we can check the present bit
    unsigned long * PTE = PTE_address(4096 * _page-no);
    int present = (*PTE & 1);
    if (present == 1)
    {
        // if present then get physical address and convert to frame number
        unsigned long physical_address = ((*PTE >> 12) << 12);
        process_mem_pool->release_frames(physical_address / 4096);
        // Turn off the present bit
        *PTE = *PTE ^ 1;
        // Reload page table to clear the TLB
        load();
    }

    Console::puts("freed page\n");
}

```

**page\_table.c: PDE\_address** : In this function, we use recursive paging to look up a PDE and return a pointer to it. This is done by creating a logical memory address that points to the page directory twice, once in the first 10 bits and again in the second 10 bits. We then return a pointer to this address that can then be de-referenced and the entry at the pointer location can then be modified.

```

unsigned long * PageTable::PDE_address(unsigned long addr)
{
    //PDE = 1023 + 1023 + given address first 10 bits plus 20 bits offset

```

```

    unsigned long PDE = (1023 << 22) | (1023 << 12) | ((addr & 1023) << 2);
    return (unsigned long *) PDE;
}

```

**page\_table.c: PTE\_address** : In this function, we use recursive paging to look up a PTE and return a pointer to it. This is done by creating a logical memory address that points to the page directory once, in the first 10 bits. We then return a pointer to this address that can then be de-referenced and the entry at the pointer location can then be modified.

```

unsigned long * PageTable::PTE_address(unsigned long addr)
{
    //PTE = 1023 + given address first 10 bits + given address second 10 bits plus 2 0 bits
    unsigned long PTE = (1023 << 22) | (((addr >> 22) & 1023) << 12) | (((addr >> 12) & 1023) << 2);
    return (unsigned long *) PTE;
}

```

**page\_table.c: handle\_fault** : Since the last MP there have been some changes in this function, the first is the implementation of recursive paging to access PTEs and PDEs. I now set PDEs and PTEs using the helper functions previously. By de-referencing the pointers the helper functions give me I'm able to manipulate the memory address in the entries whether it be setting it equal to a frame or changing the permission or valid bits. The other change is the implementation of checking whether or not a given memory address is valid in my VMpool I do this by looping thru the registered VMpools and calling is\_legitimate on each VMpool with the address I'm checking. If a get false return than we don't handle the fault.

```

void PageTable::handle_fault(REGS * _r)
{
    // create bools for bits to check
    bool bit_2 = false;
    bool bit_1 = false;
    bool bit_0 = false;

    // create masks for bits to shift
    int mask_2 = 1 << 2;
    int mask_1 = 1 << 1;
    int mask_0 = 1 << 0;

    // get error code
    int code = _r->err_code;

    // check bit_2
    if ((code & mask_2) == 0)
    {
        bit_2 = false;
    }
    else
    {
        bit_2 = true;
    }

    // check bit_1
    if ((code & mask_1) == 0)
    {
        bit_1 = false;
    }
}

```

```

}
else
{
    bit_1 = true;
}
// check bit_0
if ((code & mask_0) == 0)
{
    bit_0 = false;
}
else
{
    bit_0 = true;
}

// Check the address in the cr2 register
unsigned long attempted_address = read_cr2();

// Region check to see if address is in registered VM pool
VMPool* loop = Head;
// set pool var to false flip if address is in pool
bool in_pool = false;
// loop till there are no more pools
while (loop != nullptr)
{
    // call is legitimate to check if address belongs to pool
    in_pool = loop->is_legitimate(attempted_address);
    // break if address in pool
    if (in_pool)
    {
        break;
    }
    // got to next pool
    loop = loop->next;
}

// aborting if address is not legit
if (!in_pool)
{
    return;
}

// check page present bit
if (!bit_0)
{
    // get page directory index
    unsigned long * PDE = PDE_address(attempted_address);

    // Check whether or not the address exists in the page directory
    int mask_pde_present = 1;
    if ((*PDE & mask_pde_present) == 0)
    {
        // put the new page table page in the Page Directory
        *PDE = 4096 * process_mem_pool->get_frames(1);
    }
}

```

```

        // Kernel and user page, read and write, present PDE
        *PDE = *PDE | 7;
    }

    // Load new page into page table
    unsigned long * PTE = PTE_address(attempted_address);
    *PTE = 4096 * process_mem_pool->get_frames(1);

    // Create a new page table entry, check bit 2 to know which permissions
    if (!bit_2)
    {
        // Kernel only page, read and write, present
        *PTE = *PTE | 3;
    }
    else
    {
        // Kernel and user page, read and write, present
        *PTE = *PTE | 7;
    }
}

Console::puts("handled page fault\n");
}

```

**vm\_pool.c: VMPool Contructor** : In this function, I set up the private variables needed for the VM Pool and the 2 arrays: the allocated array and the free array. These arrays manage the memory for the VM pool when allocating in the allocate function and freeing memory in the release function.

```

VMPool::VMPool(unsigned long _base_address,
               unsigned long _size,
               ContFramePool *_frame_pool,
               PageTable *_page_table) {
    // set base address
    base_address = _base_address;
    // set size
    size = _size;
    // set frame pool ptr
    frame_pool = _frame_pool;
    // set page table ptr
    page_table = _page_table;
    // register vm pool
    page_table->register_pool(this);
    // set the allocated array to the first page
    alloc_array = (struct region *) (base_address);
    alloc_array[0].base_page = 0;
    alloc_array[0].length = 8192;
    // set the free array to the second page
    free_array = (struct region *) (base_address + 4096);
    free_array[0].base_page = 2;
    free_array[0].length = size - 8192;
    // set all the other locations in the arrays to 0
    for (int i = 1; i < 512; i++)
    {
        alloc_array[i].base_page = 0;
        alloc_array[i].length = 0;
        free_array[i].base_page = 0;
    }
}

```

```

        free_array[i].length = 0;
    }
    Console::puts("Constructed VMPool object.\n");
}

```

**vm\_pool.c: allocate** : In this function, I allocate memory in the VM pool. In order to do this I just check the amount of memory requested in the input arg `_size`. I then check if we have that amount of memory in the context of 4096 multiples in the free array as I can't allocate any memory less than 4096 since that is a page and we aren't allocating partial pages. If there is enough space in the free memory I take the memory block from the free array and put it in the allocated array.

```

unsigned long VMPool::allocate(unsigned long _size) {
    // This is always a multiple of 4096

    for (int i = 0; i < 512; i++)
    {
        // set pages and size to 0 in loop
        unsigned long pages = 0;
        unsigned long size = 0;
        // check if size is greater than 4096
        if (4096 >= _size)
        {
            // only 1 page needed
            pages = 1;
            size = 4096;
        }
        else
        {
            // Find out if an extra page is needed for anything that doesn't fit
            int remainder = _size % 4096;
            pages = _size / 4096;

            // add extra page
            if (remainder != 0)
            {
                pages++;
            }
            // set size of mem
            size = pages * 4096;
        }

        // find enough mem in free mem
        if (free_array[i].length >= size)
        {
            // find open slot in alloc mem
            for (int j = 0; j < 512; j++)
            {
                if(alloc_array[j].length == 0)
                {
                    // transfer the base page
                    alloc_array[j].base_page = free_array[i].base_page;
                    // set the memory needed
                    alloc_array[j].length = size;
                    // set the free array to the next free page and subtract the
                    // size from the length of the block
                    free_array[i].base_page = free_array[i].base_page + pages;
                    free_array[i].length = free_array[i].length - size;
                }
            }
        }
    }
}

```

```

        // return the address of the memory
        return base_address + (alloc_array[j].base_page * 4096);
    }
}

}

}

}
Console::puts("Allocated region of memory.\n");
return 0;
}

```

**vm\_pool.c: Release** : In this function, I release memory based off the address input variable passed in. I do this by looking up what the starting page number would be for the start address I think check the length for the associated base page and deallocate the pages in the contiguous block by moving the memory block to the free array and calling the free function in the page table object.

```

void VMPool::release(unsigned long _start_address) {
    // check that the start address is valid
    if (is_legitimate(_start_address))
    {
        // get the base page no
        unsigned long page_no = (_start_address - base_address) / 4096;
        // find allocation with matching base page
        for (int i = 0; i < 512; i++)
        {
            if(alloc_array[i].base_page == page_no)
            {
                // free the pages in the allocation
                for(int j = 0; j < (alloc_array[i].lenght / 4096); j++)
                {
                    unsigned long free_address = base_address
                    + ((alloc_array[i].base_page + j) * 4096);
                    page_table->free_page(free_address);
                }
                // move memory back to the free array
                for (int j = 0; j < 512; j++)
                {
                    // move mem into an free array slot not being used
                    if (free_array[j].lenght == 0)
                    {
                        // transfer the mem location and lenght
                        free_array[j].base_page = alloc_array[i].base_page;
                        free_array[j].lenght = alloc_array[i].lenght;
                        // deallocate from the allocated array
                        alloc_array[i].base_page = 0;
                        alloc_array[i].lenght = 0;
                        return;
                    }
                }
            }
        }
    }
}

```



```

    Console::puts("Released region of memory.\n");
}

```

**vm\_pool.c: is\_legitimate** : In this function, I check whether or not a given memory address passed in an input variable is legitimate. This is done by returning true if the address is in one of the allocated memory regions in the allocated array. I believe I have a bug in this function unfortunately I am out of time to fix it but don't think my logic for checking the allocated array is correct as everything passes all tests when I set it to return everything as true but when I run my code with logic uncommented I believe it produces some false negatives and returns false when it shouldn't and the program gets hung on a page fault.

```

bool VMPool::is_legitimate(unsigned long _address) {
    // handle case for first 2 pages as they should always be in mem
    unsigned long array_space = base_address + 8192;
    if ((_address >= base_address) && (_address < array_space))
    {
        return true;
    }

    // Check if the address passed in is in a valid memory region
    // for (int i = 0; i < 512; i++)
    // {
    //     if ((_address >= (base_address + (alloc_array[i].base_page * 4096)) )
    // && (_address < (base_address + (alloc_array[i].base_page * 4096)
    // + alloc_array[i].length)))
    //     {
    //         return true;
    //     }
    // }
    return true;
    Console::puts("Checked whether address is part of an allocated region.\n");
}

```

## Testing

I didn't add any testing since I believe that the given testing covers all of my code. I spent about 7 hours using the testing given to remove various bugs from my code. The first error I caught was the cast of the free\_array start location. I didn't have parenthesis around "base address + 4096" and I was casting them as a region struct. This was causing the free array to take up the same memory location as my allocated array. I then caught a logic error in my allocate function where I wasn't removing memory from my free array correctly. Lastly, there is still a logic bug in my is\_legitimate function where I'm returning a false negative if I uncomment the logic that is commented out. This is causing an infinite loop in my page fault handler as it keeps returning cause the legitimate function says the address isn't valid and we aren't supposed to handle page faults at invalid addresses.